

Handle class imbalance in #TidyTuesday climbing expedition data with tidymodels

Sep 23, 2020 ·  [rstats](#), [tidymodels](#)

Lately I've been publishing [screencasts](#) demonstrating how to use the [tidymodels](#) framework, from starting out with first modeling steps to tuning more complex models. Today's screencast walks through a detailed model analysis from beginning to end, with important feature engineering steps and several model types, using this week's [#TidyTuesday dataset](#) on Himalayan climbing expeditions. ▲

Impute missing data and handle class imbalance for Himalayan climbing expe...



Here is the code I used in the video, for those who prefer reading instead of or in addition to video.

Explore the data

Our modeling goal is to predict the probability of an Himalayan expedition member surviving or dying [based on characteristics of the person and climbing expedition from this week's #TidyTuesday dataset](#). This dataset gives us the opportunity to talk about feature engineering steps like subsampling for class imbalance (many more people survive than die) and imputing missing data (lots of expedition members are missing age, for example).

Let's start by reading in the data.

```
library(tidyverse)
members <- read_csv("https://raw.githubusercontent.com/rfordatascience/tidytuesday/master/data/2020/2020-09-22/members.csv")


members
```

```
## # A tibble: 76,519 x 21
##   expedition_id member_id peak_id peak_name year season sex age
##   <chr>          <chr>    <chr>   <chr>    <dbl> <chr>  <chr> <dbl>
## 1 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    40
## 2 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    41
## 3 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    27
## 4 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    40
## 5 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    34
## 6 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    25
## 7 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    41
## 8 AMAD78301     AMAD7830_ AMAD    Ama Dabl_ 1978 Autumn M    29
## 9 AMAD79101     AMAD7910_ AMAD    Ama Dabl_ 1979 Spring M    35
## 10 AMAD79101    AMAD7910_ AMAD    Ama Dabl_ 1979 Spring M    37
## # ... with 76,509 more rows, and 13 more variables: citizenship <chr>,
## # expedition_role <chr>, hired <lgl>, highpoint_metres <dbl>, success <lgl>,
## # solo <lgl>, oxygen_used <lgl>, died <lgl>, death_cause <chr>,
## # death_height_metres <dbl>, injured <lgl>, injury_type <chr>,
## # injury_height_metres <dbl>
```

On the video, I walk through the results of `skimr::skim()`, and notice which variables have missing data, how many unique values there are for variables like citizenship or mountain peak, and so forth.


How has the rate of expedition success and member death changed over time?

```
members %>%
  group_by(year = 10 * (year %/% 10)) %>%
  summarise(
    died = mean(died),
    success = mean(success)
  ) %>%
  pivot_longer(died:success, names_to = "outcome", values_to = "percent") %>%
  ggplot(aes(year, percent, color = outcome)) +
  geom_line(alpha = 0.7, size = 1.5) +
  scale_y_continuous(labels = scales::percent_format()) +
  labs(x = NULL, y = "% of expedition members", color = NULL)
```



Is there a relationship between the expedition member's age and success of the expedition or death? We can use the same code but just switch out `year` for `age`.

```
members %>%
  group_by(age = 10 * (age %/% 10)) %>%
  summarise(
    died = mean(died),
    success = mean(success)
  ) %>%
  pivot_longer(died:success, names_to = "outcome", values_to = "percent") %>%
  ggplot(aes(age, percent, color = outcome)) +
  geom_line(alpha = 0.7, size = 1.5) +
  scale_y_continuous(labels = scales::percent_format()) +
  labs(x = NULL, y = "% of expedition members", color = NULL)
```



Are people more likely to die on unsuccessful expeditions?

```
members %>%
  count(success, died) %>%
  group_by(success) %>%
  mutate(percent = scales::percent(n / sum(n))) %>%
  kable(
    col.names = c("Expedition success", "Died", "Number of people", "% of people"),
    align = "llrr"
  )
```

| Expedition success | Died | Number of people | % of people |
|--------------------|-------|------------------|-------------|
| FALSE | FALSE | 46452 | 98% |
| FALSE | TRUE | 868 | 2% |
| TRUE | FALSE | 28961 | 99% |
| TRUE | TRUE | 238 | 1% |

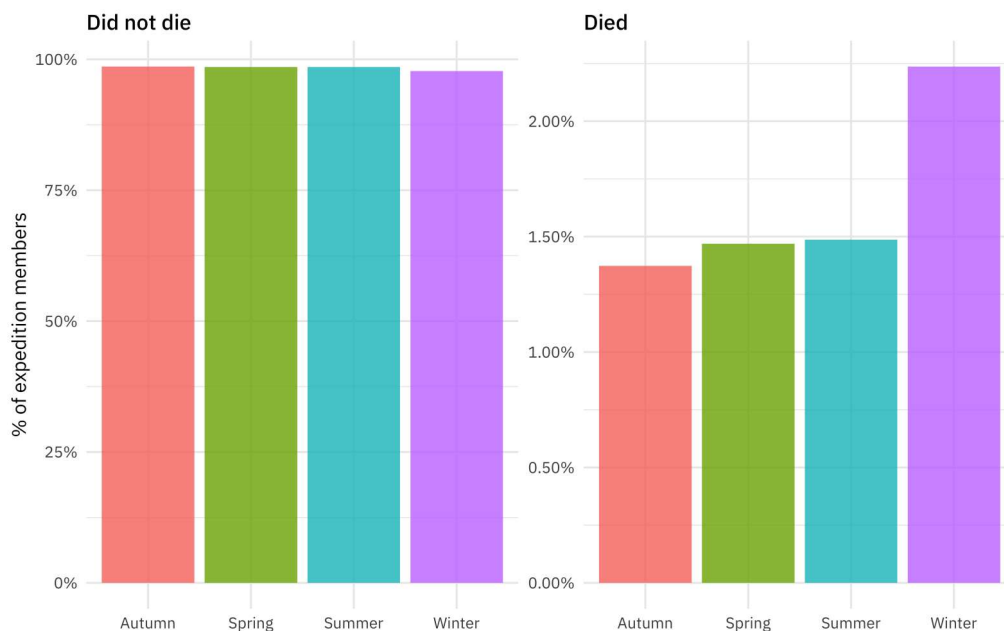
We can use a similar approach to see how different the rates of death are on different peaks in the Himalayas.

```
members %>%
  filter(!is.na(peak_name)) %>%
  mutate(peak_name = fct_lump(peak_name, prop = 0.05)) %>%
  count(peak_name, died) %>%
  group_by(peak_name) %>%
  mutate(percent = scales::percent(n / sum(n))) %>%
  kable(
    col.names = c("Peak", "Died", "Number of people", "% of people"),
    align = "llrr"
  )
```

| Peak | Died | Number of people | % of people |
|------------|-------|------------------|-------------|
| Ama Dablam | FALSE | 8374 | 100% |
| Ama Dablam | TRUE | 32 | 0% |
| Cho Oyu | FALSE | 8838 | 99% |
| Cho Oyu | TRUE | 52 | 1% |
| Everest | FALSE | 21507 | 99% |
| Everest | TRUE | 306 | 1% |
| Manaslu | FALSE | 4508 | 98% |
| Manaslu | TRUE | 85 | 2% |
| Other | FALSE | 32171 | 98% |
| Other | TRUE | 631 | 2% |

Let's make one last exploratory plot and look at seasons. How much difference is there in survival across the four seasons?

```
members %>%
  filter(season != "Unknown") %>%
  count(season, died) %>%
  group_by(season) %>%
  mutate(
    percent = n / sum(n),
    died = case_when(
      died ~ "Died",
      TRUE ~ "Did not die"
    )
  ) %>%
  ggplot(aes(season, percent, fill = season)) +
  geom_col(alpha = 0.8, position = "dodge", show.legend = FALSE) +
  scale_y_continuous(labels = scales::percent_format()) +
  facet_wrap(~died, scales = "free") +
  labs(x = NULL, y = "% of expedition members")
```



There are lots more great examples of #TidyTuesday EDA out there to explore on [Twitter](#)! Let's now create the dataset that we'll use for modeling by filtering on some of the variables and transforming some variables to a few factors. There are still lots of NA values for age but we are going to *impute* those.

```
members_df <- members %>%
  filter(season != "Unknown", !is.na(sex), !is.na(citizenship)) %>%
  select(peak_id, year, season, sex, age, citizenship, hired, success, died) %>%
  mutate(died = case_when(
    died ~ "died",
    TRUE ~ "survived"
  )) %>%
  mutate_if(is.character, factor) %>%
  mutate_if(is.logical, as.integer)

members_df
```

```
## # A tibble: 76,507 x 9
##   peak_id year season sex   age citizenship hired success died
##   <dbl> <dbl> <fct> <fct> <dbl> <fct>      <int> <int> <fct>
## 1 AMAD  1978 Autumn M     40 France      0      0 survived
## 2 AMAD  1978 Autumn M     41 France      0      0 survived
## 3 AMAD  1978 Autumn M     27 France      0      0 survived
## 4 AMAD  1978 Autumn M     40 France      0      0 survived
## 5 AMAD  1978 Autumn M     34 France      0      0 survived
## 6 AMAD  1978 Autumn M     25 France      0      0 survived
## 7 AMAD  1978 Autumn M     41 France      0      0 survived
## 8 AMAD  1978 Autumn M     29 France      0      0 survived
## 9 AMAD  1979 Spring M     35 USA        0      0 survived
## 10 AMAD 1979 Spring M     37 W Germany 0      1 survived
## # ... with 76,497 more rows
```

Build a model

We can start by loading the tidymodels metapackage, and splitting our data into training and testing sets.

```
library(tidymodels)

set.seed(123)
members_split <- initial_split(members_df, strata = died)
members_train <- training(members_split)
members_test <- testing(members_split)
```

We are going to use [resampling](#) to evaluate model performance, so let's get those resampled sets ready.

```
set.seed(123)
members_folds <- vfold_cv(members_train, strata = died)
members_folds
```

```
## # 10-fold cross-validation using stratification
## # A tibble: 10 x 2
##   splits      id
##   <list>      <chr>
## 1 <split [51.6K/5.7K]> Fold01
## 2 <split [51.6K/5.7K]> Fold02
## 3 <split [51.6K/5.7K]> Fold03
## 4 <split [51.6K/5.7K]> Fold04
## 5 <split [51.6K/5.7K]> Fold05
## 6 <split [51.6K/5.7K]> Fold06
## 7 <split [51.6K/5.7K]> Fold07
## 8 <split [51.6K/5.7K]> Fold08
## 9 <split [51.6K/5.7K]> Fold09
## 10 <split [51.6K/5.7K]> Fold10
```

Next we build a recipe for data preprocessing.

- First, we must tell the `recipe()` what our model is going to be (using a formula here) and what our training data is.
- Next, we impute the missing values for `age` using the median age in the training data set. There are more complex [steps available for imputation](#), but we'll stick with a straightforward option here.
- Next, we use `step_other()` to collapse categorical levels for peak and citizenship. Before this step, there were hundreds of values in each variable.
- After this, we can create indicator variables for the non-numeric, categorical values, except for the outcome `died` which we need to keep as a factor.
- Finally, there are many more people who survived their expedition than who died (thankfully) so [we will use step_smote\(\)](#) to balance the classes.

The object `members_rec` is a recipe that has **not** been trained on data yet (for example, which categorical levels should be collapsed has not been calculated).

```
library(themis)

members_rec <- recipe(died ~ ., data = members_train) %>%
  step_medianimpute(age) %>%
  step_other(peak_id, citizenship) %>%
  step_dummy(all_nominal(), -died) %>%
  step_smote(died)

members_rec
```

```
## Data Recipe
##
## Inputs:
##
##   role #variables
##   outcome      1
##   predictor      8
##
## Operations:
##
## Median Imputation for age
## Collapsing factor levels for peak_id, citizenship
## Dummy variables from all_nominal(), -died
## SMOTE based on died
```

We're going to use this recipe in a `workflow()` so we don't need to stress a lot about whether to `prep()` or not. If you want to explore the what the recipe is doing to your data, you can first `prep()` the recipe to estimate the parameters needed for each step and then `bake(new_data = NULL)` to pull out the training data with those steps applied.

Let's compare *two* different models, a logistic regression model and a random forest model; these are the same two models I used [in the post on the Palmer penguins](#). We start by creating the model specifications.

```
glm_spec <- logistic_reg() %>%
  set_engine("glm")

glm_spec
```

```
## Logistic Regression Model Specification (classification)
##
## Computational engine: glm
```

```
rf_spec <- rand_forest(trees = 1000) %>%
  set_mode("classification") %>%
  set_engine("ranger")

rf_spec
```

```
## Random Forest Model Specification (classification)
##
## Main Arguments:
##   trees = 1000
##
## Computational engine: ranger
```

Next let's start putting together a tidymodels `workflow()`, a helper object to help manage modeling pipelines with pieces that fit together like Lego blocks. Notice that there is no model yet: **Model: None**.

```
members_wf <- workflow() %>%
  add_recipe(members_rec)

members_wf
```

```
## == Workflow ==
## Preprocessor: Recipe
## Model: None
##
## -- Preprocessor --
## 4 Recipe Steps
##
## • step_medianimpute()
## • step_other()
## • step_dummy()
## • step_smote()
```

Now we can add a model, and the fit to each of the resamples. First, we can fit the logistic regression model. Let's set a non-default metric set so we can add sensitivity and specificity.

```
members_metrics <- metric_set(roc_auc, accuracy, sensitivity, specificity)

doParallel::registerDoParallel()
glm_rs <- members_wf %>%
  add_model(glm_spec) %>%
  fit_resamples(
    resamples = members_folds,
    metrics = members_metrics,
    control = control_resamples(save_pred = TRUE)
  )

glm_rs
```

```
## # Resampling results
## # 10-fold cross-validation using stratification
## # A tibble: 10 x 5
##   splits          id .metrics      .notes      .predictions
##   <list>         <chr> <list>      <list>      <list>
## 1 <split [51.6K/5.7K... Fold01 <tibble [4 x 3... <tibble [0 x ... <tibble [5,739 x 5...
## 2 <split [51.6K/5.7K... Fold02 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 3 <split [51.6K/5.7K... Fold03 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 4 <split [51.6K/5.7K... Fold04 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 5 <split [51.6K/5.7K... Fold05 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 6 <split [51.6K/5.7K... Fold06 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 7 <split [51.6K/5.7K... Fold07 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 8 <split [51.6K/5.7K... Fold08 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 9 <split [51.6K/5.7K... Fold09 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 10 <split [51.6K/5.7K... Fold10 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
```

Second, we can fit the random forest model.

```
rf_rs <- members_wf %>%
  add_model(rf_spec) %>%
  fit_resamples(
    resamples = members_folds,
    metrics = members_metrics,
    control = control_resamples(save_pred = TRUE)
  )

rf_rs
```

```
## # Resampling results
## # 10-fold cross-validation using stratification
## # A tibble: 10 x 5
##   splits          id .metrics      .notes      .predictions
##   <list>         <chr> <list>      <list>      <list>
## 1 <split [51.6K/5.7K... Fold01 <tibble [4 x 3... <tibble [0 x ... <tibble [5,739 x 5...
## 2 <split [51.6K/5.7K... Fold02 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 3 <split [51.6K/5.7K... Fold03 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 4 <split [51.6K/5.7K... Fold04 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 5 <split [51.6K/5.7K... Fold05 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 6 <split [51.6K/5.7K... Fold06 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 7 <split [51.6K/5.7K... Fold07 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 8 <split [51.6K/5.7K... Fold08 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 9 <split [51.6K/5.7K... Fold09 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
## 10 <split [51.6K/5.7K... Fold10 <tibble [4 x 3... <tibble [0 x ... <tibble [5,738 x 5...
```

We have fit each of our candidate models to our resampled training set!

Evaluate model

Now let's check out how we did.

```
collect_metrics(glm_rs)
```

```
## # A tibble: 4 x 5
##   .metric .estimator mean    n std_err
##   <chr>   <chr>      <dbl> <int>  <dbl>
## 1 accuracy binary     0.619   10 0.00230
## 2 roc_auc  binary     0.705   10 0.00721
## 3 sens     binary     0.678   10 0.0139
## 4 spec     binary     0.619   10 0.00243
```

Well, this is middling but at least mostly consistent for the positive and negative classes. The function `collect_metrics()` extracts and formats the `.metrics` column from resampling results like the ones we have here.


```
collect_metrics(rf_rs)
```

```
## # A tibble: 4 x 5
##   .metric .estimator mean     n std_err
##   <chr>   <chr>   <dbl> <int>  <dbl>
## 1 accuracy binary    0.972   10 0.000514
## 2 roc_auc  binary    0.746   10 0.00936
## 3 sens     binary    0.164   10 0.0125
## 4 spec     binary    0.984   10 0.000499
```

The accuracy is great but that sensitivity... YIKES! The random forest model has not done a great job of learning how to recognize *both* classes, even with our oversampling strategy. Let's dig deeper into how these models are doing to see this more. For example, how are they predicting the two classes?

```
glm_rs %>%
  conf_mat_resampled()
```

```
## # A tibble: 4 x 3
##   Prediction Truth     Freq
##   <fct>         <fct>   <dbl>
## 1 died        died     55.5
## 2 died        survived 2157.
## 3 survived    died      26.5
## 4 survived    survived 3499.
```


```
rf_rs %>%
  conf_mat_resampled()
```

```
## # A tibble: 4 x 3
##   Prediction Truth     Freq
##   <fct>         <fct>   <dbl>
## 1 died        died     13.5
## 2 died        survived  89.9
## 3 survived    died     68.5
## 4 survived    survived 5566.
```

The random forest model is quite bad at identifying which expedition members died, while the logistic regression model does about the same for both classes.

We can also make an ROC curve.

```
glm_rs %>%
  collect_predictions() %>%
  group_by(id) %>%
  roc_curve(died, .pred_died) %>%
  ggplot(aes(1 - specificity, sensitivity, color = id)) +
  geom_abline(lty = 2, color = "gray80", size = 1.5) +
  geom_path(show.legend = FALSE, alpha = 0.6, size = 1.2) +
  coord_equal()
```



It is finally time for us to return to the testing set. Notice that we have not used the testing set yet during this whole analysis; to compare and assess models we used resamples of the training set. Let's *fit* one more time to the training data and *evaluate* on the testing data using the function `last_fit()`.

```
members_final <- members_wf %>%
  add_model(glm_spec) %>%
  last_fit(members_split)

members_final
```

```
## # Resampling results
## # Manual resampling
## # A tibble: 1 x 6
##   splits      id      .metrics  .notes  .predictions  .workflow
##   <list>    <chr>    <list>    <list>    <list>        <list>
## 1 <split [57.4K... train/test... <tibble [2 ... <tibble [0... <tibble [19,126... <workflo...
```

The metrics and predictions here are on the *testing* data.

```
collect_metrics(members_final)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy binary      0.619
## 2 roc_auc  binary      0.689
```

```
collect_predictions(members_final) %>%
  conf_mat(died, .pred_class)
```

```
##           Truth
## Prediction  died survived
## died        196    7204
## survived    90    11636
```

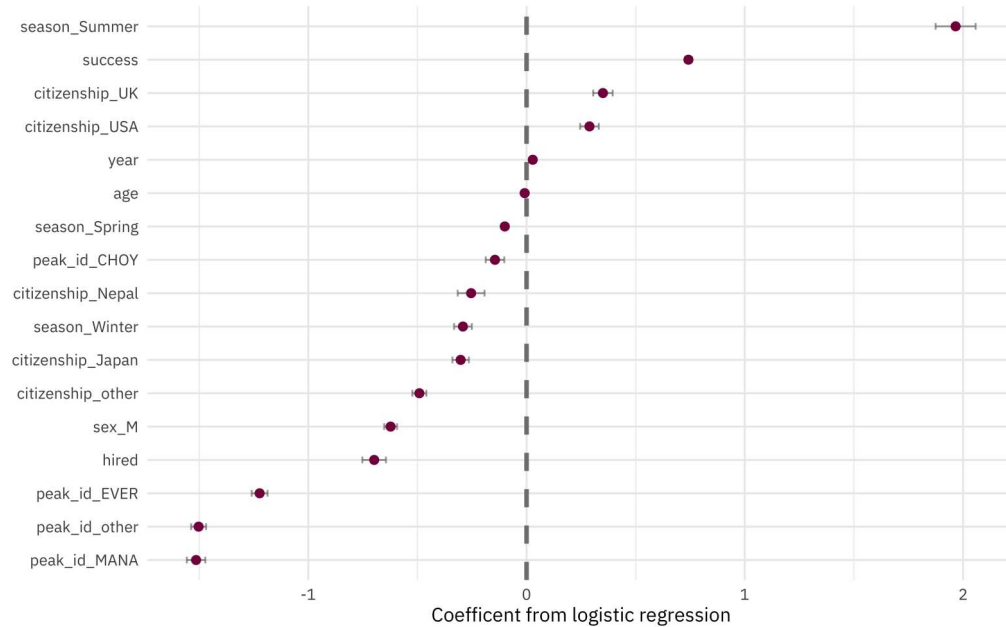
The coefficients (which we can get out using `tidy()`) have been estimated using the *training* data. If we use `exponentiate = TRUE`, we have odds ratios.

```
members_final %>%
  pull(.workflow) %>%
  pluck(1) %>%
  tidy(exponentiate = TRUE) %>%
  arrange(estimate) %>%
  kable(digits = 3)
```

| term | estimate | std.error | statistic | p.value |
|-------------------|----------|-----------|-----------|---------|
| (Intercept) | 0.000 | 0.944 | -57.309 | 0.000 |
| peak_id_MANA | 0.220 | 0.042 | -35.769 | 0.000 |
| peak_id_other | 0.223 | 0.034 | -43.635 | 0.000 |
| peak_id_EVER | 0.294 | 0.036 | -33.641 | 0.000 |
| hired | 0.497 | 0.054 | -12.928 | 0.000 |
| sex_M | 0.536 | 0.029 | -21.230 | 0.000 |
| citizenship_other | 0.612 | 0.032 | -15.299 | 0.000 |
| citizenship_Japan | 0.739 | 0.038 | -7.995 | 0.000 |
| season_Winter | 0.747 | 0.041 | -7.180 | 0.000 |
| citizenship_Nepal | 0.776 | 0.062 | -4.128 | 0.000 |
| peak_id_CHOY | 0.865 | 0.043 | -3.404 | 0.001 |
| season_Spring | 0.905 | 0.016 | -6.335 | 0.000 |
| age | 0.991 | 0.001 | -12.129 | 0.000 |
| year | 1.029 | 0.000 | 59.745 | 0.000 |
| citizenship_USA | 1.334 | 0.043 | 6.759 | 0.000 |
| citizenship_UK | 1.419 | 0.045 | 7.858 | 0.000 |
| success | 2.099 | 0.016 | 46.404 | 0.000 |
| season_Summer | 7.142 | 0.092 | 21.433 | 0.000 |

We can also visualize these results.


```
members_final %>%
  pull(.workflow) %>%
  pluck(1) %>%
  tidy() %>%
  filter(term != "(Intercept)") %>%
  ggplot(aes(estimate, fct_reorder(term, estimate))) +
  geom_vline(xintercept = 0, color = "gray50", lty = 2, size = 1.2) +
  geom_errorbar(aes(
    xmin = estimate - std.error,
    xmax = estimate + std.error
  ),
  width = .2, color = "gray50", alpha = 0.7
) +
  geom_point(size = 2, color = "#85144B") +
  labs(y = NULL, x = "Coefficient from logistic regression")
```



- The features with coefficients on the positive side (like climbing in summer, being on a successful expedition, or being from the UK or US) are associated with surviving.
- The features with coefficients on the negative side (like climbing specific peaks including Everest, being one of the hired members of an expedition, or being a man) are associated with dying.

Remember that we have to interpret model coefficients like this in light of the predictive accuracy of our model, which was somewhat middling; there are more factors at play in who survives these expeditions than what we have accounted for in the model directly. Also note that we see evidence in this model for how dangerous it is to be a native Sherpa climber in Nepal, hired as an expedition member, as [pointed out in this week's #TidyTuesday README](#).

[rstats](#) [tidymodels](#)



Julia Silge

Data Scientist & Software Engineer

I'm an author, international keynote speaker, and real-world practitioner focusing on data analysis and machine learning practice. I love making beautiful charts and communicating about technical topics with diverse audiences.



Related

- [Introducing our new book, Tidy Modeling with R](#)
- [Train and analyze many models for #TidyTuesday crop yields](#)
- [Build a #TidyTuesday predictive text model for The Last Airbender](#)
- [Get started with tidymodels and #TidyTuesday Palmer penguins](#)
- [Bagging with tidymodels and #TidyTuesday astronaut missions](#)

